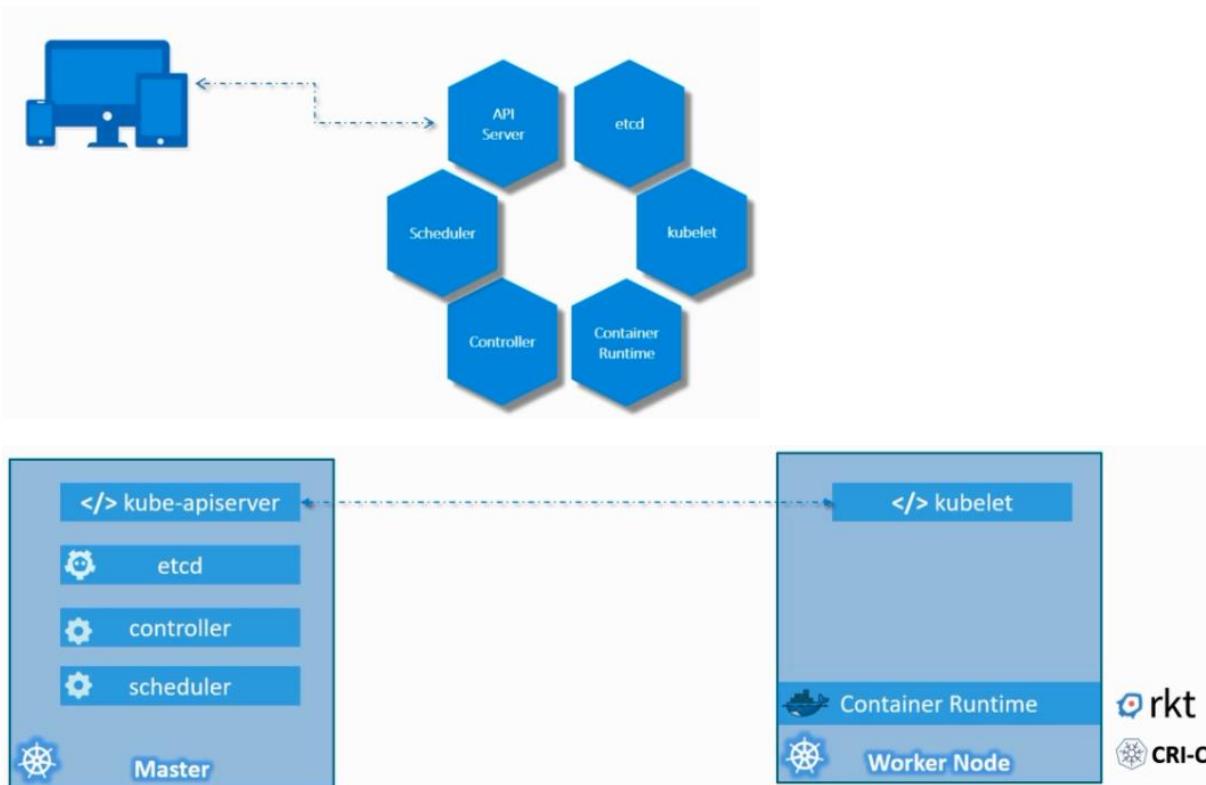


Core Concepts - 13%

1. Architecture



kubectl get nodes

kubectl cluster-info

2. Pod:

```
kubectl run --generator=run-pod/v1 nginx-pod --image=nginx  
kubectl run --generator=run-pod/v1 nginx --image=nginx --port=8080 -l tier=db --env="DNS_DOMAIN=cluster" --dry-run -o yaml  
kubectl run hazelcast --image=nginx --env="DNS_DOMAIN=cluster" --env="POD_NAMESPACE=default"  
kubectl run hazelcast --image=nginx -l app=name,wf=ov  
kubectl run --generator=run-pod/v1 nginx --image=nginx --port=8080 --command echo hi  
kubectl create -f *.yaml  
kubectl edit pod pod-name      - edit already running pod // Need to verify  
kubectl apply -f *.yaml  
kubectl get pod <podname> -o wide  
kubectl delete pod <podname>  
kubectl describe pod <podname>  
kubectl get pod podname -o yaml --export  
kubectl set image pod/nginx nginx=nginx:1.7.1  
kubectl exec podname -it bash or kubectl exec podname -it -- /bin/sh    # get inside the pod  
kubectl exec -it nginx - env  # to display env variables  
Pod deleted automatically when its completed ** use --restart=Never
```

3. ReplicationController/ ReplicaSet

```
kubectl run replicontrolr --generator=run/v1 --image=redis --replicas=2 --dry-run -o yaml    // replicationcontroller  
To generate file for replicaset -> generate deployment file and edit the file  
kubectl create -f *.yaml  
kubectl get replicationcontroller  
kubectl get replicasesets
```

```
kubectl delete replicaset <rp-name>
kubectl describe replicaset
kubectl edit replicaset replicaset-name      - edit already running replicaset
```

4. Scale

1. update replicas number in the file

```
kubectl replace -f file.yaml
```

2. kubectl scale --replicas=6 -f file.yaml

or

```
kubectl scale --replicas=6 replicaset replicaset-name
```

```
kubectl scale --replicas=3 -f filename.yaml
```

```
kubectl scale --current-replicas=1 --replicas=3 deployment/mydeploy
```

```
kubectl scale --replicas=3 rc/rc1 rc/rc2 rc/rc3    # Multiple replication controller update
```

<<this option wont update the replicas in the file>>

```
kubectl edit rs rs-name
```

5. Deployment

```
kubectl create -f *.yaml --record
```

```
kubectl run nginx --image nginx
```

```
kubectl delete deployment dep-name
```

```
kubectl run --generator=deployment/v1beta1 nginx --image=nginx
```

Or

```
kubectl create deployment --image=nginx nginx
```

```
kubectl run --generator=deployment/v1beta1 nginx --image=nginx --dry-run -o yaml
```

Or

```
kubectl create deployment --image=nginx nginx --dry-run -o yaml
```

```
kubectl run --generator=deployment/v1beta1 nginx --image=nginx --dry-run --replicas=4 -o yaml
```

kubectl create deployment does not have a --replicas option. You could first create it and then scale it using the kubectl scale command.

```
kubectl run --generator=deployment/v1beta1 nginx --image=nginx --dry-run --replicas=4 -o yaml > nginx-deployment.yaml
```

Service:

```
kubectl create service clusterip ngservice --tcp=80:80 --dry-run -o yaml | after creating change the selector to required pod
```

```
kubectl create service nodeport nginx --tcp=80:8000 --node-port=30080 --dry-run -o yaml
```

```
kubectl expose deployment nginx --type=NodePort --port=80 --target-port=8000 --name=nginx-serv --dry-run -o yaml
```

```
kubectl expose deployment nginx --port=80 --target-port=8000
```

// need to check with docs

namespace:

```
kubectl get all --all-namespaces
```

```
kubectl exec -it app cat /log/app.log --namespace=elastic-stack
```

```
kubectl create namespace mynamespace
```

```
kubectl run nginx --image=nginx --n mynamespace
```

ResourceQuota: kubectl create quota myrq --hard(cpu=1, memory=1G, pods=2) --dry-run -o yaml

Multipod container: 10%

1. Many containers running in a single pod.

Create a pod with 2 container and run commands in the second container

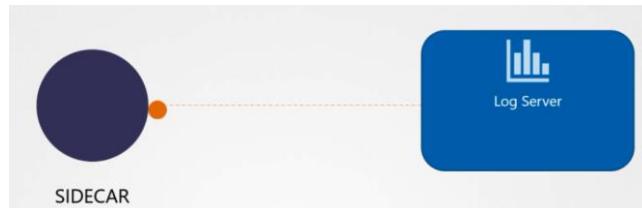
```
kubectl exec -it pod-name -c container-name-2 -- /bin/sh
```

It shares same network,volumes .

Ex: App+Log agent

Patterns : Side car , Adapter, ambassador

Side car:



Before sending log to the central server, we like to convert to the common format. For this we use **adaptor** pattern



Ambassador pattern:

We have multiple kind of DB depends on the dev,test,prod environment.
we have container that proxy the appropriate DB
Though we have 3 patterns, implementing them is always the same.

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp
  labels:
    name: simple-webapp
spec:
  containers:
  - name: simple-webapp
    image: simple-webapp
    ports:
      - containerPort: 8080
  - name: log-agent
    image: log-agent
```

POD DESIGN – 20%

Labels and Selector

When we have multiple pods, inorder to classify , separate -> Labels are used

```
Kubectl get pods --selector app=App1
```

```
Kubectl get all --selector app=app1,tier=db
```

```
Kubectl get pod --show-labels // to display all labels
```

```
Kubectl get pod -L app
```

// capital L for only specifying Key .

```
Kubectl get pod -l app=ov // small l for specify with key and value
```

```
kubectl label po nginx1 nginx2 nginx3 app- // to remove app label from the pods
```

```
kubectl label po nginx2 app=v2 --overwrite // Change the app label
```

```
kubectl annotate po nginx1 nginx2 nginx3 description='my description'
```

```
kubectl annotate po nginx1 nginx2 nginx3 description-
```

```
kubectl autoscale deploy nginx --min=5 --max=10 --cpu-percent=80 --dry-run -o yaml
```

Annotations:

Annotations under metadata is to mention build number, contact details, build version etc

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: simple-webapp
  labels:
    app: Appl
    function: Front-end
  annotations:
    buildversion: 1.34
spec:
  replicas: 3
  selector:
    matchLabels:
      app: Appl
  template:
    metadata:
      labels:
        app: Appl
        function: Front-end
    spec:
      containers:
        - name: simple-webapp
          image: simple-webapp
```

Rolling Updates & Rollbacks in deployments

Rollout

Kubectl rollout status deployment/my-app-deployment

Kubectl rollout history deployment/my-app-deployment

When first creating deployment , following steps happen

1. New rollout

2. New revision

Update – change any config in the file and apply it

After changing deplymnt file , >>> kubectl apply -f deploy-file

Kubectl set image deployment/mydeploy nginx=nginx:1.9.1

Strategy -> Recreate and Rolling update / can be viewd from the below command

Kubectl describe deployment mydeploy

The screenshot displays two columns of deployment events. The left column, labeled 'Recreate', shows events where the deployment controller scales the replica set from 0 to 5 and back to 0. The right column, labeled 'RollingUpdate', shows events where the deployment controller scales the replica set from 5 to 2, 2 to 3, 3 to 4, 4 to 5, and finally back to 2, while maintaining at least one replica available.

| Type | Reason | Age | From | Message |
|--------|-------------------|-----|-----------------------|--|
| Normal | ScalingReplicaSet | 1m | deployment-controller | Scaled up replica set myapp-deployment-6795844b58 to 5 |
| Normal | ScalingReplicaSet | 1m | deployment-controller | Scaled down replica set myapp-deployment-6795844b58 to 0 |
| Normal | ScalingReplicaSet | 56s | deployment-controller | Scaled up replica set myapp-deployment-54c7d0ccc to 5 |

| Type | Reason | Age | From | Message |
|--------|-------------------|-----|-----------------------|--|
| Normal | ScalingReplicaSet | 3w | deployment-controller | Scaled up replica set myapp-deployment-67c749c58c to 5 |
| Normal | ScalingReplicaSet | 3s | deployment-controller | Scaled up replica set myapp-deployment-7d57dbdb8d to 2 |
| Normal | ScalingReplicaSet | 3s | deployment-controller | Scaled down replica set myapp-deployment-67c749c58c to 4 |
| Normal | ScalingReplicaSet | 3s | deployment-controller | Scaled up replica set myapp-deployment-7d57dbdb8d to 3 |
| Normal | ScalingReplicaSet | 8s | deployment-controller | Scaled down replica set myapp-deployment-67c749c58c to 3 |
| Normal | ScalingReplicaSet | 8s | deployment-controller | Scaled up replica set myapp-deployment-7d57dbdb8d to 4 |
| Normal | ScalingReplicaSet | 8s | deployment-controller | Scaled down replica set myapp-deployment-67c749c58c to 2 |
| Normal | ScalingReplicaSet | 8s | deployment-controller | Scaled up replica set myapp-deployment-7d57dbdb8d to 5 |
| Normal | ScalingReplicaSet | 8s | deployment-controller | Scaled down replica set myapp-deployment-67c749c58c to 1 |

Rollback

To undo the change run :

kubectl rollout undo deployment/mydeploy

kubectl rollout undo deploy nginx --to-revision=2

kubectl rollout history deploy nginx --revision=3

kubectl rollout pause deploy nginx // to pause the rollout

kubectl rollout resume deploy nginx

Create

```
> kubectl create -f deployment-definition.yml
```

Get

```
> kubectl get deployments
```

Update

```
> kubectl apply -f deployment-definition.yml
```

Status

```
> kubectl rollout status deployment/myapp-deployment
```

```
> kubectl rollout history deployment/myapp-deployment
```

Rollback

```
> kubectl rollout undo deployment/myapp-deployment
```

Jobs

pods are designed to run all the time by k8s. So if it has a small process/less time helping task, even it is finished the task, it'll go to completed state. Then replicaset restart the pod to make it always running. To stop restarting we have to set restartPolicy: Never

```
kubectl create job busybox --image=busybox -- /bin/sh -c 'echo hello;sleep 30;echo world'
```

I Parallelism

job-definition.yaml

```
apiVersion: batch/v1
kind: Job
metadata:
  name: random-error-job
spec:
  completions: 3
  parallelism: 3
  template:
    spec:
      containers:
        - name: random-error
          image: kodekloud/random-error
  restartPolicy: Never
```

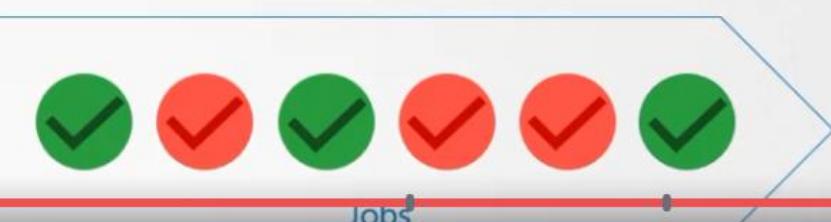
▶ kubectl create -f job-definition.yaml

▶ kubectl get jobs

| NAME | DESIRED | SUCCESSFUL | AGE |
|------------------|---------|------------|-----|
| random-error-job | 3 | 2 | 38s |

▶ kubectl get pods

| NAME | READY | STATUS | RESTARTS |
|-----------------------|-------|-----------|----------|
| random-exit-job-kmttt | 0/1 | Completed | 0 |
| random-exit-job-sdsrf | 0/1 | Error | 0 |
| random-exit-job-wwqbn | 0/1 | Completed | 0 |
| random-exit-job-fkhfn | 0/1 | Error | 0 |
| random-exit-job-fvf5t | 0/1 | Error | 0 |
| random-exit-job-nmghp | 0/1 | Completed | 0 |

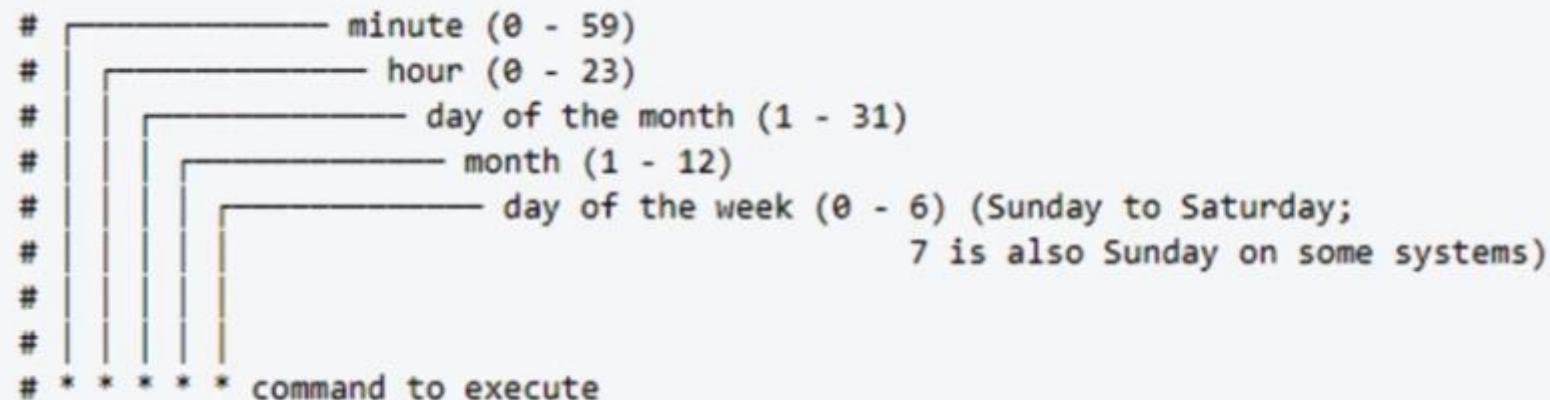


Kubectl get jobs

Kubectl delete job job-name // deleting the job always delete the pod that was created by job

Like replicas -> we have completions for job to make sure it will run the mentioned number of time. By default jobs create one after another. In order to create more jobs simultaneously use parallelism.

CronJob : it is a job which can be scheduled



Wikipedia

```
kubectl run --generator=job/v1 --image=kodekloud/throw-dice throw-dice-job --restart=Never  
kubectl run --generator=cronjob/v1beta1 --image=kodekloud/throw-dice throw-dice-cron-job --restart=Never -o yaml --dry-run --  
schedule="30 21 * * *" > cjb.yaml  
or kubectl run --generator=cronjob/v2alpha1  
once job created, if we edit the job file -> its not working for some properties.  
You have to type completions and parallelism  
backoffLimit: 25 # This is so the job does not quit before it succeeds  
kubectl create cronjob busybox --image=busybox --schedule="*/1 * * * *" -- /bin/sh -c 'date; echo Hello from the  
Kubernetes cluster'
```

Create CronJob

```
kubectl create -f cron-job-definition.yaml
```

```
kubectl get cronjob
```

| NAME | SCHEDULE | SUSPEND | ACTIVE |
|--------------------|-------------|---------|--------|
| reporting-cron-job | */1 * * * * | False | 0 |

```
cron-job-definition.yaml
```

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: reporting-cron-job
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      completions: 3
      parallelism: 3
      template:
        spec:
          containers:
            - name: reporting-tool
              image: reporting-tool
  restartPolicy: Never
```

Configurations- 18%

Command and arguments

Docker commands:

Docker containers only lives as long as process running.

Once process inside container stops, the container go to exited state

Inside Docker file -> CMD[] it runs when the container starts

```
FROM Ubuntu
```

```
CMD sleep 5
```

```
▶ docker run ubuntu-sleeper sleep 10
```

Command at Startup: sleep 10

```
FROM Ubuntu
```

```
ENTRYPOINT ["sleep"]
```

```
▶ docker run ubuntu-sleeper 10
```

Command at Startup: sleep 10

Command line arg over rides the command in the file. But if we use Entry point it appends the command and argument from the terminal
If user not sending 10 as arguments, then it ll fail at the startup. So inoder to give default value use CMD .

The screenshot shows a terminal window with a purple border containing three lines of Dockerfile code:

```
FROM Ubuntu
ENTRYPOINT ["sleep"]
CMD ["5"]
```

To the right of the terminal window is a black terminal window showing the output of the command:

```
▶ docker run ubuntu-sleeper
sleep: missing operand
Try 'sleep --help' for more information.
```

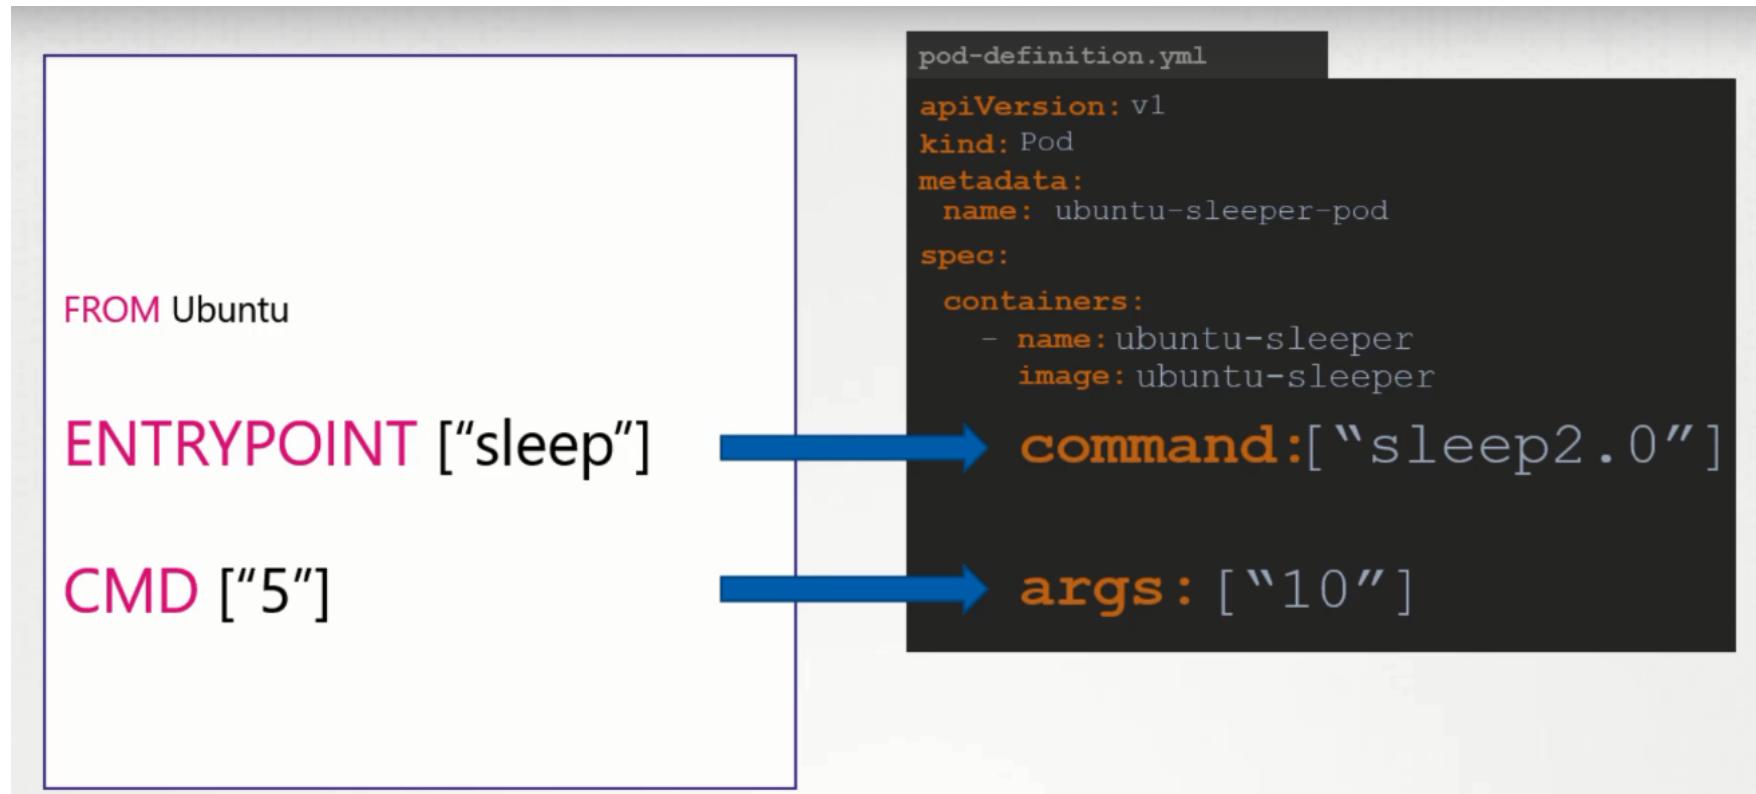
Below the terminal windows, two green boxes contain the text:

Command at Startup: sleep 5

Command at Startup: sleep 10

Use –entry-point when running docker command to over ride the entry point command.

K8s Command and arguments



Over ride the Entry point command in the docker file → use command in k8s

Override the default CMD in docker file → use args in k8s

ENV Variables in Kuber

```
▶ docker run -e APP_COLOR=pink simple-webapp-color
```

pod-definition.yaml

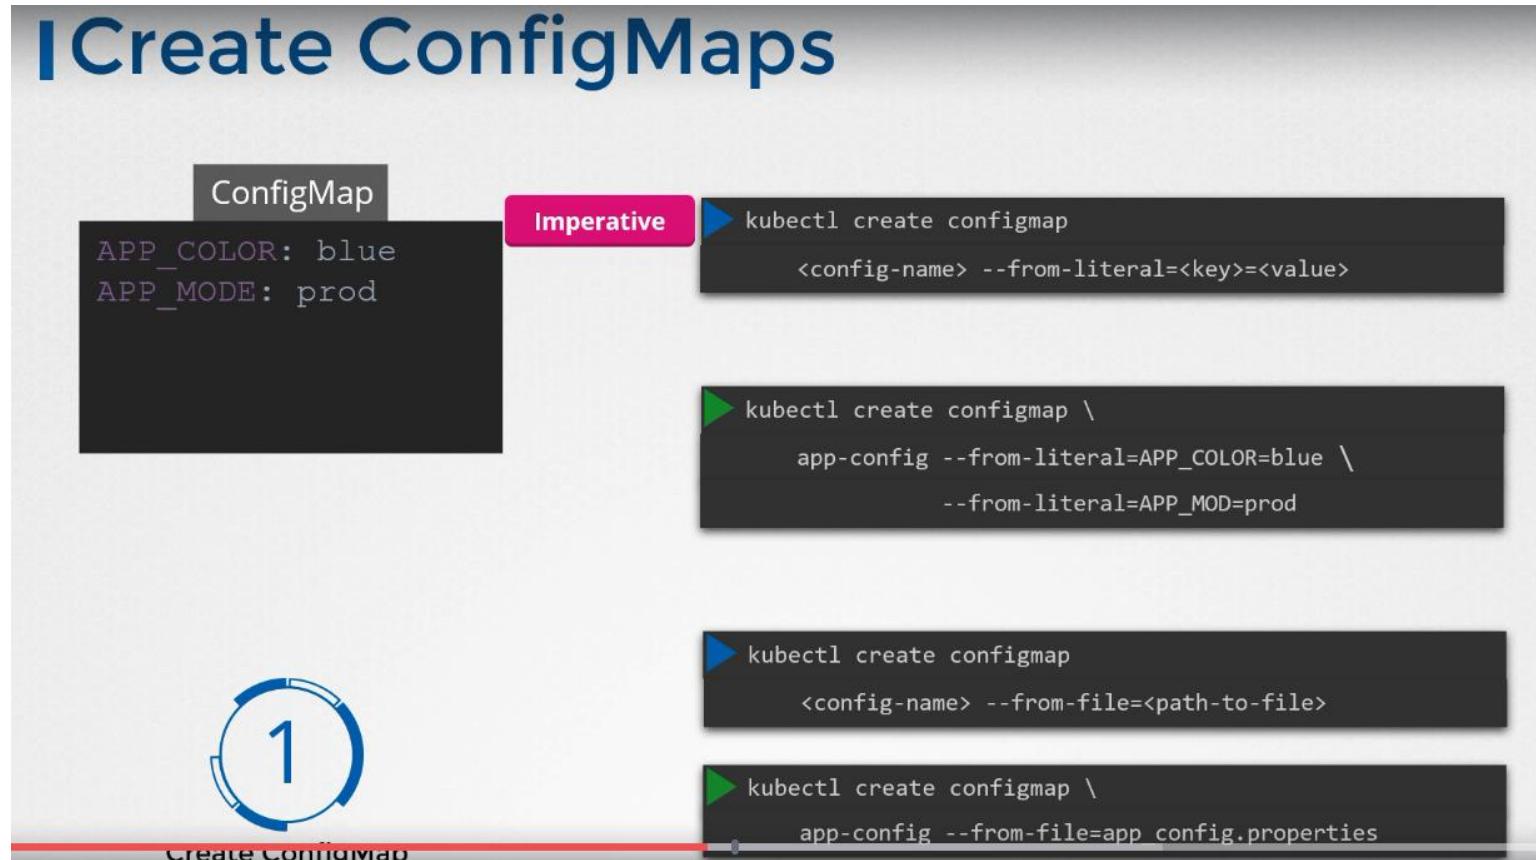
```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
spec:
  containers:
    - name: simple-webapp-color
      image: simple-webapp-color
      ports:
        - containerPort: 8080
      env:
        - name: APP_COLOR
          value: pink
```

```
kubectl run --generator=run-pod/v1 nginx --image=nginx --port=8080 -l tier=db --env="DNS_DOMAIN=cluster" --dry-run -o yaml
```

ConfigMaps

When we have lot of pod files and difficult to manage environmental variable details. So Env variables can be taken out of the pod definition file and put it in a separate file called configmaps. Config maps are key value pair in k8s.

```
Kubectl create configmap app-config --from-literal=Name=ov --from-literal=Hname=karthi
```



Create ConfigMaps

ConfigMap Declarative ➤ `kubectl create -f`

```
APP_COLOR: blue  
APP_MODE: prod
```

```
config-map.yaml
```

```
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: app-config  
data:  
  APP_COLOR: blue  
  APP_MODE: prod
```

1 Create ConfigMap ➤ `kubectl create -f config-map.yaml`

```
app-config
```

```
APP_COLOR: blue  
APP_MODE: prod
```

```
mysql-config
```

```
port: 3306  
max_allowed_packet: 128M
```

```
redis-config
```

```
port: 6379  
rdb-compression: yes
```

Kubectl get configmaps -> cmd

ConfigMap in Pods

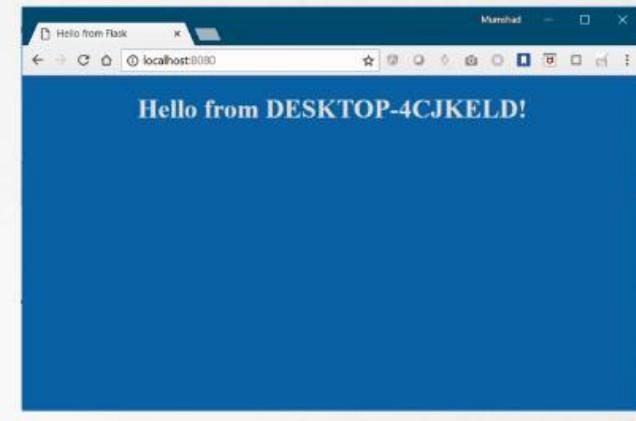
pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
  labels:
    name: simple-webapp-color
spec:
  containers:
  - name: simple-webapp-color
    image: simple-webapp-color
    ports:
    - containerPort: 8080
  envFrom:
  - configMapRef:
      name: app-config
```

```
kubectl create -f pod-definition.yaml
```

config-map.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  APP_COLOR: blue
  APP_MODE: prod
```

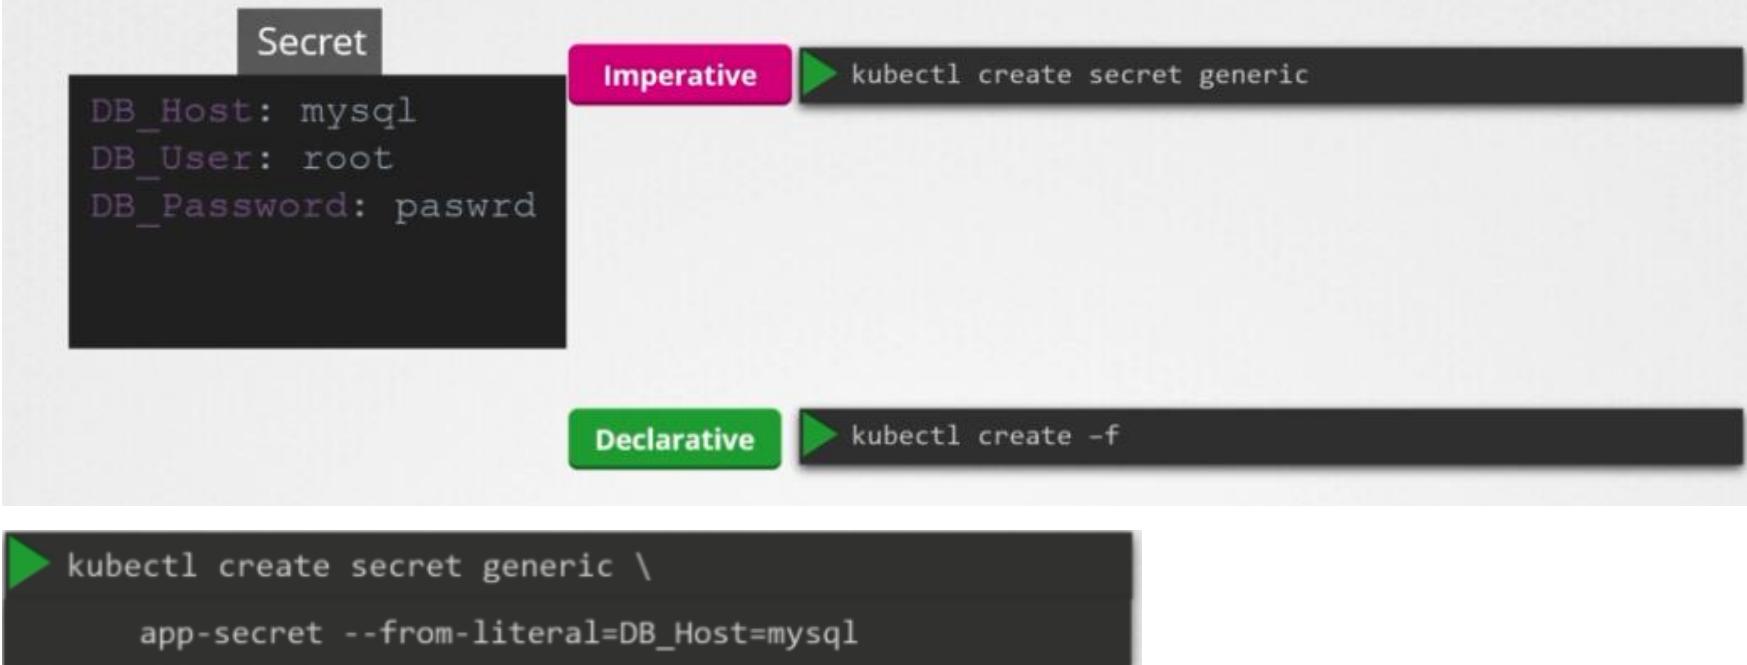


Secrets:

Secrets are used to store the sensitive informations, like password



Create Secrets



Kubectl create secret generic app-secret --from-literal=DB_HOST=mysql

Imperative

```
▶ kubectl create secret generic  
    <secret-name> --from-literal=<key>=<value>
```

```
▶ kubectl create secret generic \  
    app-secret --from-literal=DB_Host=mysql      \  
              --from-literal=DB_User=root  
              --from-literal=DB_Password=paswrd
```

```
▶ kubectl create secret generic  
    <secret-name> --from-file=<path-to-file>
```

```
▶ kubectl create secret generic \  
    app-secret --from-file=app_secret.properties
```

The diagram illustrates the process of creating a Secret using declarative methods. It consists of three main sections:

- Secret**: A tab labeled "Secret" is selected.
- Declarative**: A tab labeled "Declarative" is selected.
- kubectl create -f**: A command-line interface (CLI) prompt with the command "kubectl create -f".

Create Secret: A circular icon with a blue border and a white center containing the number "1". Below it is the text "Create Secret".

secret-data.yaml: A file named "secret-data.yaml" containing the following YAML configuration:

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
data:
  DB_Host: bXlzcWw=
  DB_User: cm9vdA==
  DB_Password: cGFzd3Jk
```

kubectl create -f secret-data.yaml: A CLI prompt with the command "kubectl create -f secret-data.yaml".

echo -n 'mysql' | base64: A CLI prompt with the command "echo -n 'mysql' | base64". The output is "bXlzcWw=". This row has a light gray background.

echo -n 'root' | base64: A CLI prompt with the command "echo -n 'root' | base64". The output is "cm9vdA==". This row has a light gray background.

echo -n 'paswrd' | base64: A CLI prompt with the command "echo -n 'paswrd' | base64". The output is "cGFzd3Jk". This row has a light gray background.

```
▶ kubectl get secrets
NAME          TYPE        DATA   AGE
app-secret    Opaque      3      10m
default-token-mvtkv  kubernetes.io/service-account-token  3      2h

▶ kubectl describe secrets
Name:         app-secret
Namespace:    default
Labels:       <none>
Annotations: <none>

Type:        Opaque

Data
=====
DB_Host:    10 bytes
DB_Password: 6 bytes
DB_User:    4 bytes

▶ kubectl get secret app-secret -o yaml
apiVersion: v1
data:
  DB_Host: bXlzcWw=
  DB_Password: cGFzd3Jk
  DB_User: cm9vdA==
kind: Secret
metadata:
  creationTimestamp: 2018-10-18T10:01:12Z
  labels:
    name: app-secret
    name: app-secret
    namespace: default
  uid: be96e989-d2bc-11e8-a545-080027931072
type: Opaque
```

Decode the hash value

```
▶ echo -n 'bXlzcWw=' | base64 --decode
mysql
```

1 Secrets in Pods

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
  labels:
    name: simple-webapp-color
spec:
  containers:
    - name: simple-webapp-color
      image: simple-webapp-color
      ports:
        - containerPort: 8080
  envFrom:
    - secretRef:
        name: app-secret
```

secret-data.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
data:
  DB_Host: bXlzcWw=
  DB_User: cm9vdA==
  DB_Password: cGFzd3Jk
```



Inject into Pod

```
envFrom:  
  - secretRef:  
      name: app-config
```

ENV

SINGLE ENV

```
env:  
  - name: DB_Password  
    valueFrom:  
      secretKeyRef:  
        name: app-secret  
        key: DB_Password
```

```
volumes:  
  - name: app-secret-volume  
    secret:  
      secretName: app-secret
```

VOLUME

Docker Security:

For a host -> root user || For a Docker container → root user

Security contexts

Pod level:

I Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  securityContext:
    runAsUser: 1000

  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
```

Container level:

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
      securityContext:
        runAsUser: 1000
```

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
      securityContext:
        runAsUser: 1000
      capabilities:
        add: ["MAC_ADMIN"]
```



Note: Capabilities are only supported at the container level and not at the POD level

```
kubectl exec ubuntu-sleeper whoami
```

```
kubectl exec -it ubuntu-sleeper -- date -s '19 APR 2012 11:14:00'
```

Service Accounts:



```
▶ kubectl create serviceaccount dashboard-sa
serviceaccount "dashboard-sa" created
```



```
▶ kubectl get serviceaccount
  NAME          SECRETS   AGE
  default       1         218d
  dashboard-sa  1         4d
```



```
▶ kubectl describe serviceaccount dashboard-sa
  Name:           dashboard-sa
  Namespace:      default
  Labels:         <none>
  Annotations:   <none>
  Image pull secrets: <none>
  Mountable secrets: dashboard-sa-token-kbbdm
  Tokens:         dashboard-sa-token-kbbdm
  Events:        <none>
```

```
▶ kubectl describe secret dashboard-sa-token-kbbdm
```

```
Name:           dashboard-sa-token-kbbdm
Namespace:      default
Labels:         <none>
Type:          kubernetes.io/service-account-token
```

```
Data
```

```
====
```

```
ca.crt:    1025 bytes
namespace:  7 bytes
token:
eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcmsgdGVzL
3NlcnpY2VhY2NvdW50Iiwia3ViZXJuZXRob3cv9zZXJ2aWNLYWNib3V
```

```
▶ kubectl exec -it my-kubernetes-dashboard ls /var/run/secrets/kubernetes.io/serviceaccount
```

```
ca.crt  namespace  token
```

```
kubectl get serviceaccount
```

| NAME | SECRETS | AGE |
|--------------|---------|------|
| default | 1 | 218d |
| dashboard-sa | 1 | 4d |

pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: my-kubernetes-dashboard
spec:
  containers:
    - name: my-kubernetes-dashboard
      image: my-kubernetes-dashboard
  serviceAccount: dashboard-sa
```

```
pod-definition.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-kubernetes-dashboard
spec:
  containers:
    - name: my-kubernetes-dashboard
      image: my-kubernetes-dashboard
  automountServiceAccountToken: false
```

You can replace the POD type with Deployment, then add your **serviceAccountName: default** under

```
1 | template:
2 | spec:
3 |   serviceAccountName: default
```

Resources Requirement:



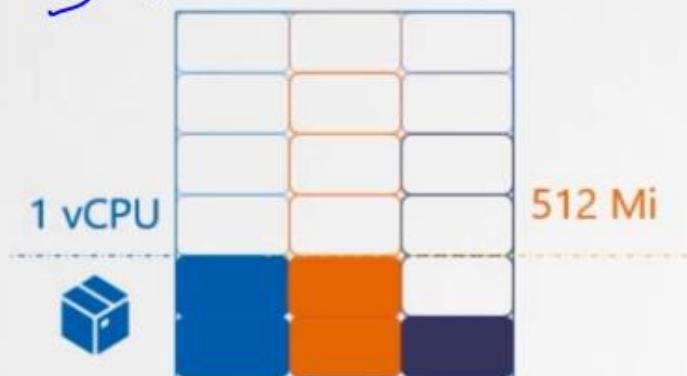
Default:



We can also set limit for resource usages in the pod

I Resource Limits

Default



pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
  labels:
    name: simple-webapp-color
spec:
  containers:
  - name: simple-webapp-color
    image: simple-webapp-color
    ports:
    - containerPort: 8080
  resources:
    requests:
      memory: "1Gi"
      cpu: 1
    limits:
      memory: "2Gi"
      cpu: 2
```

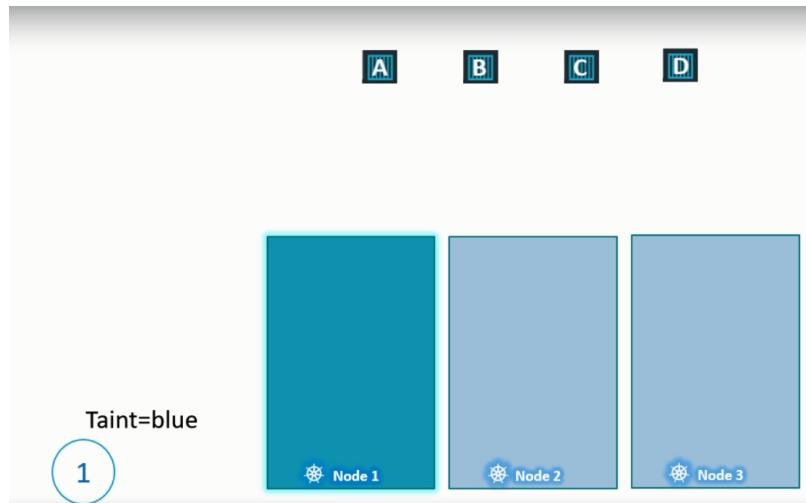
Container can use the more memory resources than its limit, But not CpU

Taints and Tolerants



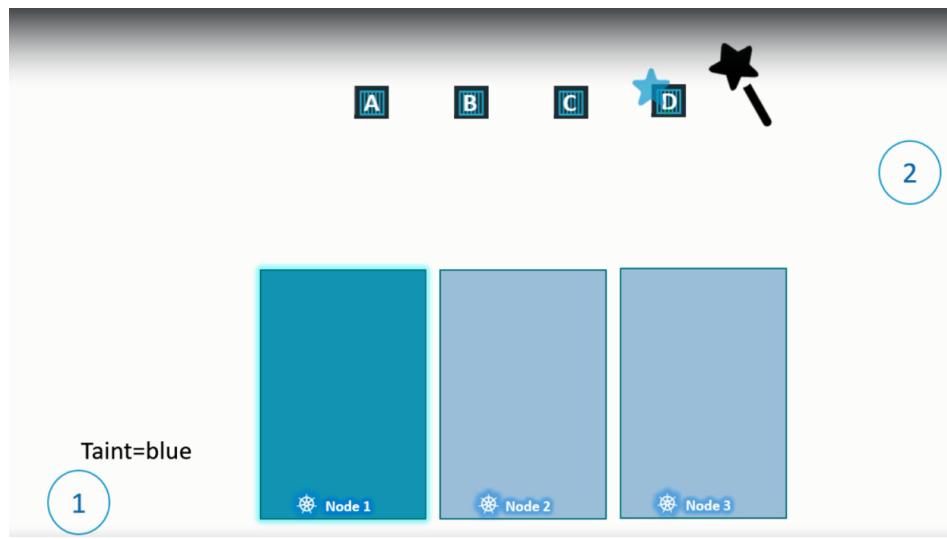
Imp 1. Taints on Node 2. Tolerant level

For some specific use cases , we nee dto deploy the pod on specific Nodes.



No unwanted pods ll be placed on Node1

Taints on set on Nodes , Tolerance set on Pods



🔧 Taints - Node

```
kubectl taint nodes node-name key=value:taint-effect
```

NoSchedule | PreferNoSchedule | NoExecute

What happens to PODs
that do not tolerate this taint?

```
kubectl taint nodes node1 app=blue:NoSchedule
```

★ Tolerations - PODs

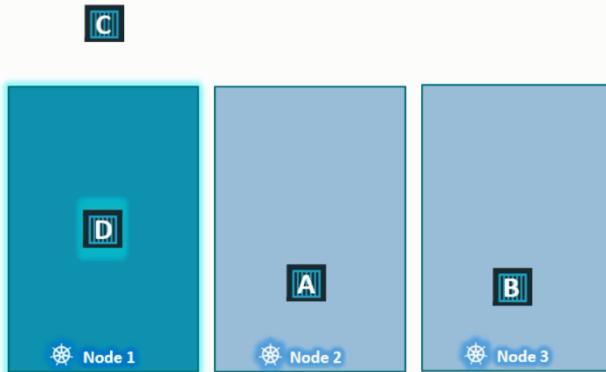
```
pod-definition.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
  - name: nginx-container
    image: nginx
  tolerations:
  - key: "app"
    operator: "Equal"
    value: "blue"
    effect: "NoSchedule"
```

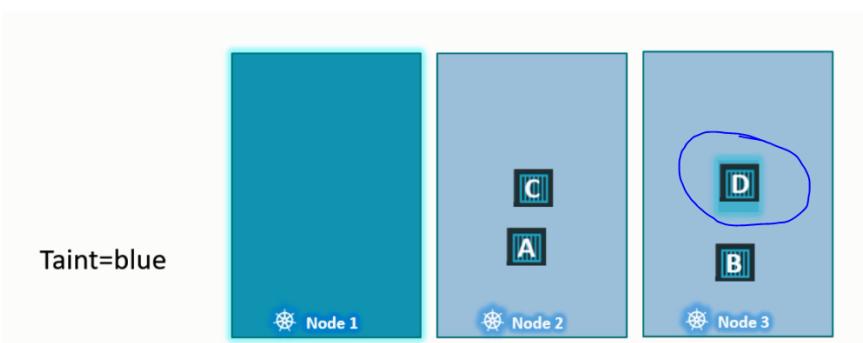
Initialy pod C deployed in the Node1, before adding taint to node1.

Pod C from Node1 was evicted because it cannot tolerate the newly added taint

Taint - NoExecute



No Gurantee that pod D always deployed in Node1. It could be deployed in Node3.



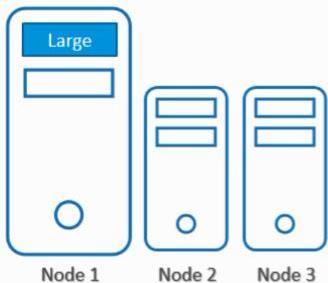
BY default Master node is tainted to avoid deploying pod.

```
kubectl describe node kubemaster | grep Taint  
Taints:           node-role.kubernetes.io/master:NoSchedule
```

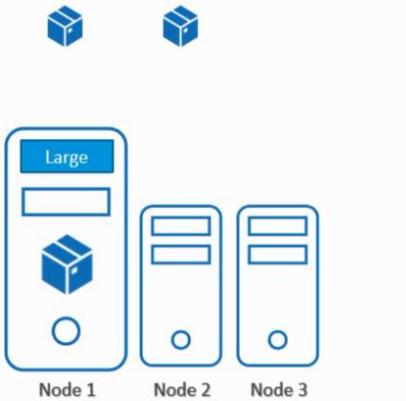
Node Selector

Label Nodes

```
▶ kubectl label nodes <node-name> <label-key>=<label-value>  
▶ kubectl label nodes node-1 size=Large
```



Node Selector

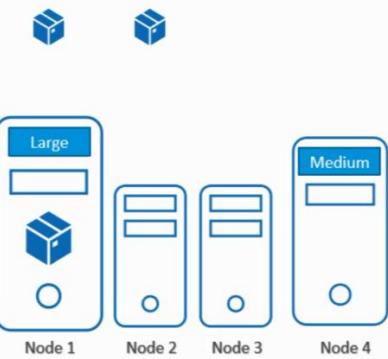


```
pod-definition.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: data-processor
      image: data-processor
  nodeSelector:
    size: Large
```

```
▶ kubectl create -f pod-definition.yml
```

Some cases cannot be achieved by Node selector

Node Selector - Limitations



- Large OR Medium?
- NOT Small

Advanced operators like – OR, AND, IN, NOTIN, EXISTS

Node affinity – To add extra operators so that pod is deployed on desired node

Node Affinity

pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: data-processor
      image: data-processor
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: size
                operator: NotIn
                values:
                  - Small
```

Many operators are available , ex – In, NotIn, Exists

Node Affinity Types

Available:

`requiredDuringSchedulingIgnoredDuringExecution`
`preferredDuringSchedulingIgnoredDuringExecution`

Planned:

`requiredDuringSchedulingRequiredDuringExecution`

Node Affinity Types

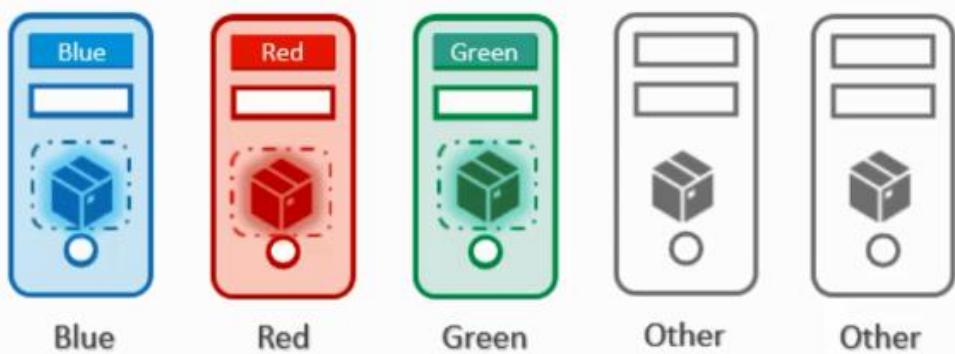
Available:

`requiredDuringSchedulingIgnoredDuringExecution`
`preferredDuringSchedulingIgnoredDuringExecution`

| | DuringScheduling | DuringExecution |
|--------|------------------|-----------------|
| Type 1 | Required | Ignored |
| Type 2 | Preferred | Ignored |

A pod is running in a Node, in future the label of the Node is changed, in that case pod is continued to be run on the same Node.

Taints/Tolerations and Node Affinity



Observability- 18%

Have to search liveness probe in search get the template

Readiness Probe

pod-definition.yaml



```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp
  labels:
    name: simple-webapp
spec:
  containers:
  - name: simple-webapp
    image: simple-webapp
    ports:
    - containerPort: 8080
  readinessProbe:
    httpGet:
      path: /api/ready
      port: 8080
```

HTTP Test - /api/ready

I Readiness Probe

```
readinessProbe:  
  httpGet:  
    path: /api/ready  
    port: 8080  
  
  initialDelaySeconds: 10  
  
  periodSeconds: 5  
  
  failureThreshold: 8
```

```
readinessProbe:  
  tcpSocket:  
    port: 3306
```

```
readinessProbe:  
  exec:  
    command:  
      - cat  
      - /app/is_ready
```

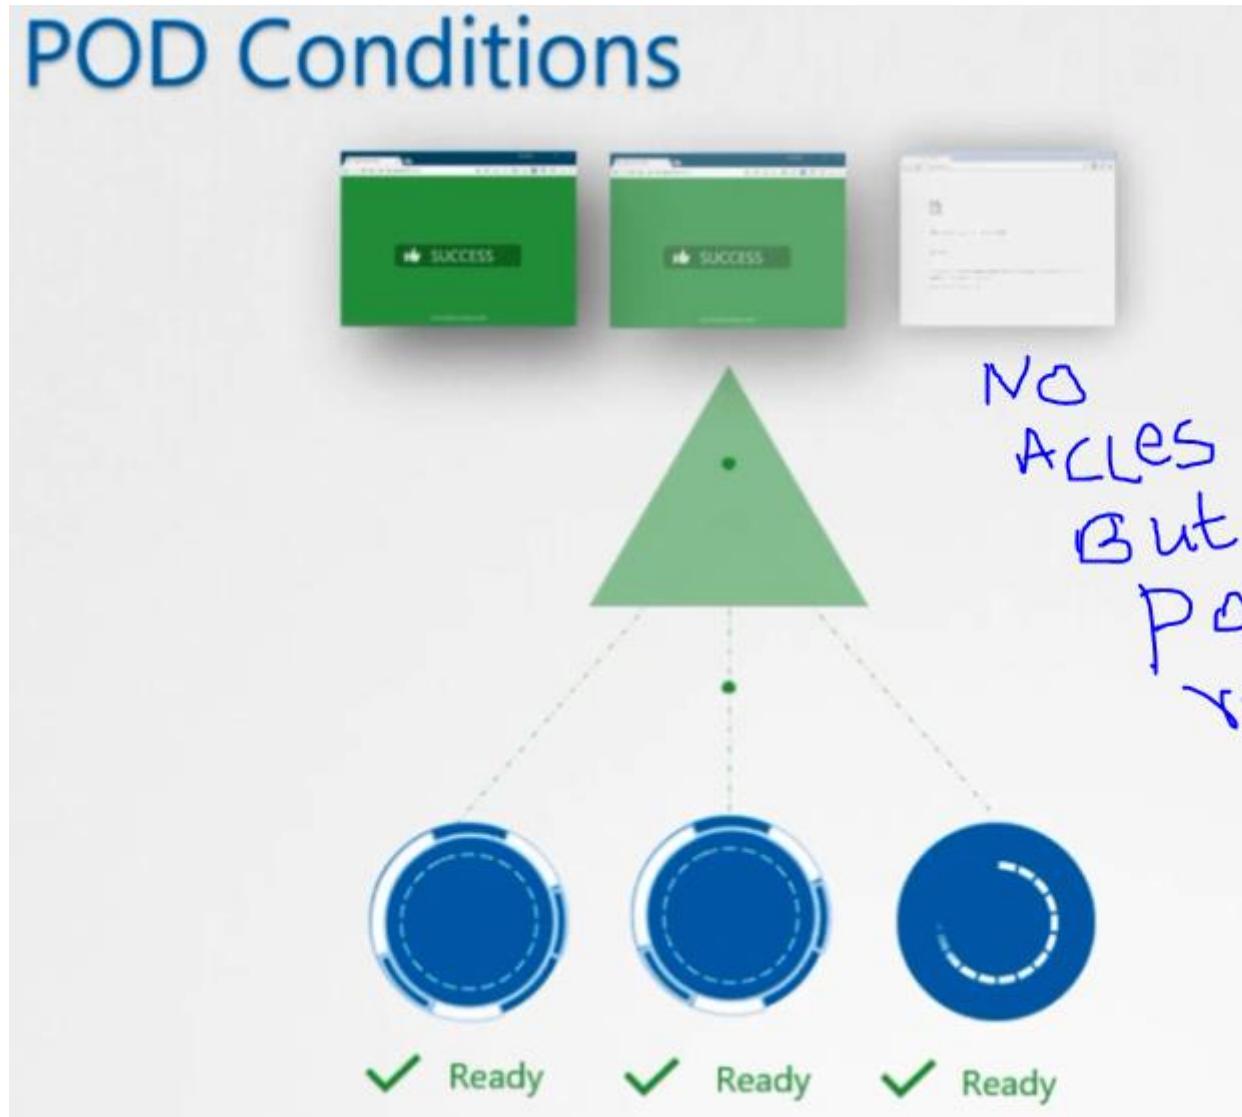
HTTP Test - /api/ready

TCP Test - 3306

Exec Command

Default threshold is 3 attempts, after that it declare it as fail. But we can increase threshold by failureThreshold option

POD Conditions



Liveness Probes



HTTP Test - /api/healthy



TCP Test - 3306



Exec Command

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp
  labels:
    name: simple-webapp
spec:
  containers:
  - name: simple-webapp
    image: simple-webapp
    ports:
    - containerPort: 8080
      livenessProbe:
        httpGet:
          path: /api/healthy
          port: 8080
```

Similar to readinessProbe

I Logs - Docker

```
▶ docker run -d kodekloud/event-simulator
```

```
▶ docker logs -f ecf
```

```
2018-10-06 15:57:15,937 - root - INFO - USER1 logged in
2018-10-06 15:57:16,943 - root - INFO - USER2 logged out
2018-10-06 15:57:17,944 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:18,951 - root - INFO - USER3 is viewing page3
2018-10-06 15:57:19,954 - root - INFO - USER4 is viewing page1
2018-10-06 15:57:20,955 - root - INFO - USER2 logged out
2018-10-06 15:57:21,956 - root - INFO - USER1 logged in
2018-10-06 15:57:22,957 - root - INFO - USER3 is viewing page2
2018-10-06 15:57:23,959 - root - INFO - USER1 logged out
2018-10-06 15:57:24,959 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:25,961 - root - INFO - USER1 logged in
2018-10-06 15:57:26,965 - root - INFO - USER4 is viewing page3
2018-10-06 15:57:27,965 - root - INFO - USER4 is viewing page3
2018-10-06 15:57:28,967 - root - INFO - USER2 is viewing page1
2018-10-06 15:57:29,967 - root - INFO - USER3 logged out
```

I Logs - Kubernetes

```
▶ kubectl create -f event-simulator.yaml
```

```
▶ kubectl logs -f event-simulator-pod
```

```
2018-10-06 15:57:15,937 - root - INFO - USER1 logged in
2018-10-06 15:57:16,943 - root - INFO - USER2 logged out
2018-10-06 15:57:17,944 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:18,951 - root - INFO - USER3 is viewing page3
2018-10-06 15:57:19,954 - root - INFO - USER4 is viewing page1
2018-10-06 15:57:20,955 - root - INFO - USER2 logged out
2018-10-06 15:57:21,956 - root - INFO - USER1 logged in
```

event-simulator.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: event-simulator-pod
spec:
  containers:
    - name: event-simulator
      image: kodekloud/event-simulator
```

multi container pod

```
▶ kubectl logs -f event-simulator-pod event-simulator
2018-10-06 15:57:15,937 - root - INFO - USER1 logged in
2018-10-06 15:57:16,943 - root - INFO - USER2 logged out
2018-10-06 15:57:17,944 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:18,951 - root - INFO - USER3 is viewing page3
2018-10-06 15:57:19,954 - root - INFO - USER4 is viewing page1
2018-10-06 15:57:20,955 - root - INFO - USER2 logged out
```

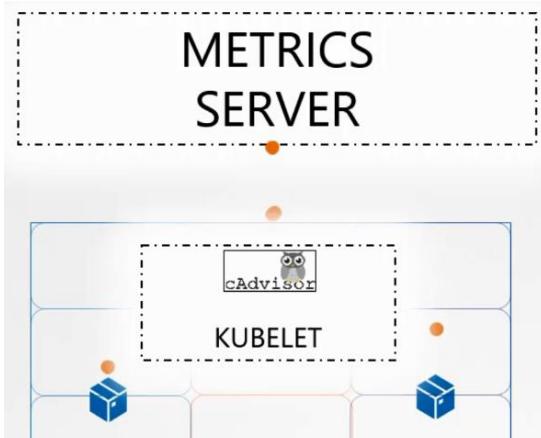
event-simulator.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: event-simulator-pod
spec:
  containers:
    - name: event-simulator
      image: kodekloud/event-simulator
    - name: image-processor
      image: some-image-processor
```

Monitoring: Node level, Pod level – collect metrics of Cpu,Ram, memory usages and give proper analytics

Many 3rd party open source solutions are available. Now we look only Metrics Server

Metrics-Server: In memory only



After installing Metrics server in k8s cluster run the below command to see the metrics

```
▶ kubectl top node
```

| NAME | CPU(cores) | CPU% | MEMORY(bytes) | MEMORY% |
|------------|------------|------|---------------|---------|
| kubemaster | 166m | 8% | 1337Mi | 70% |
| kubenode1 | 36m | 1% | 1046Mi | 55% |
| kubenode2 | 39m | 1% | 1048Mi | 55% |

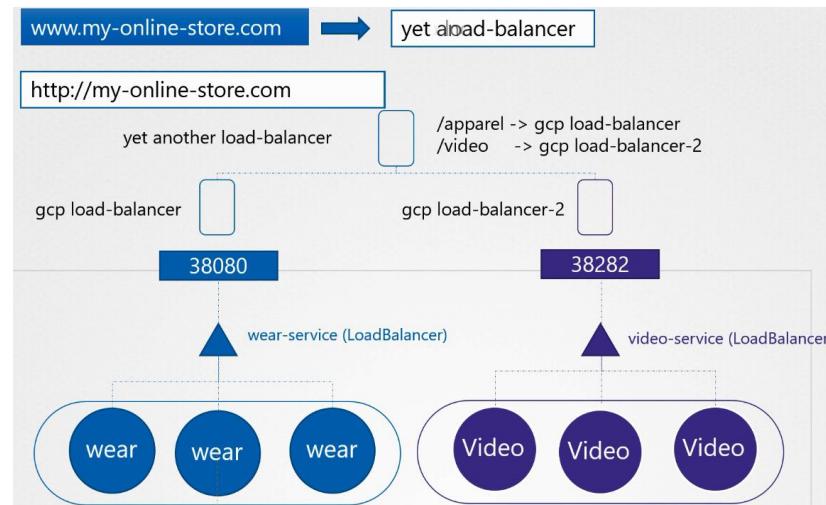
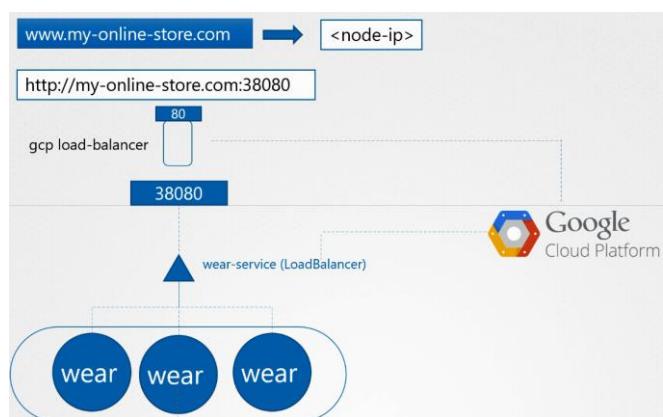
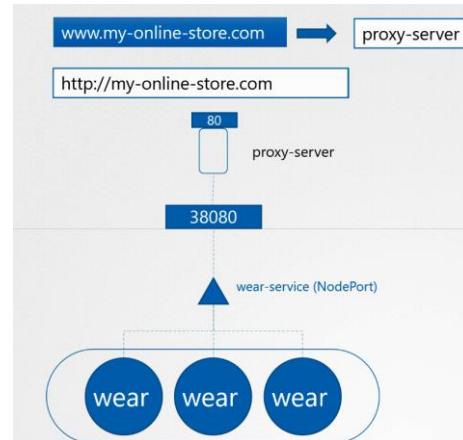
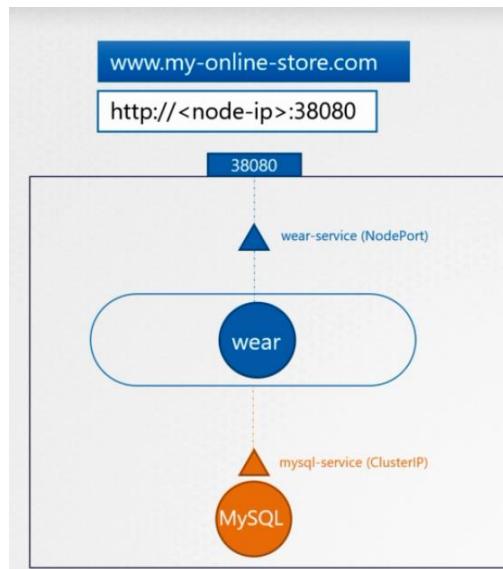
Network and Services – 13%

```
service-definition.yml
```

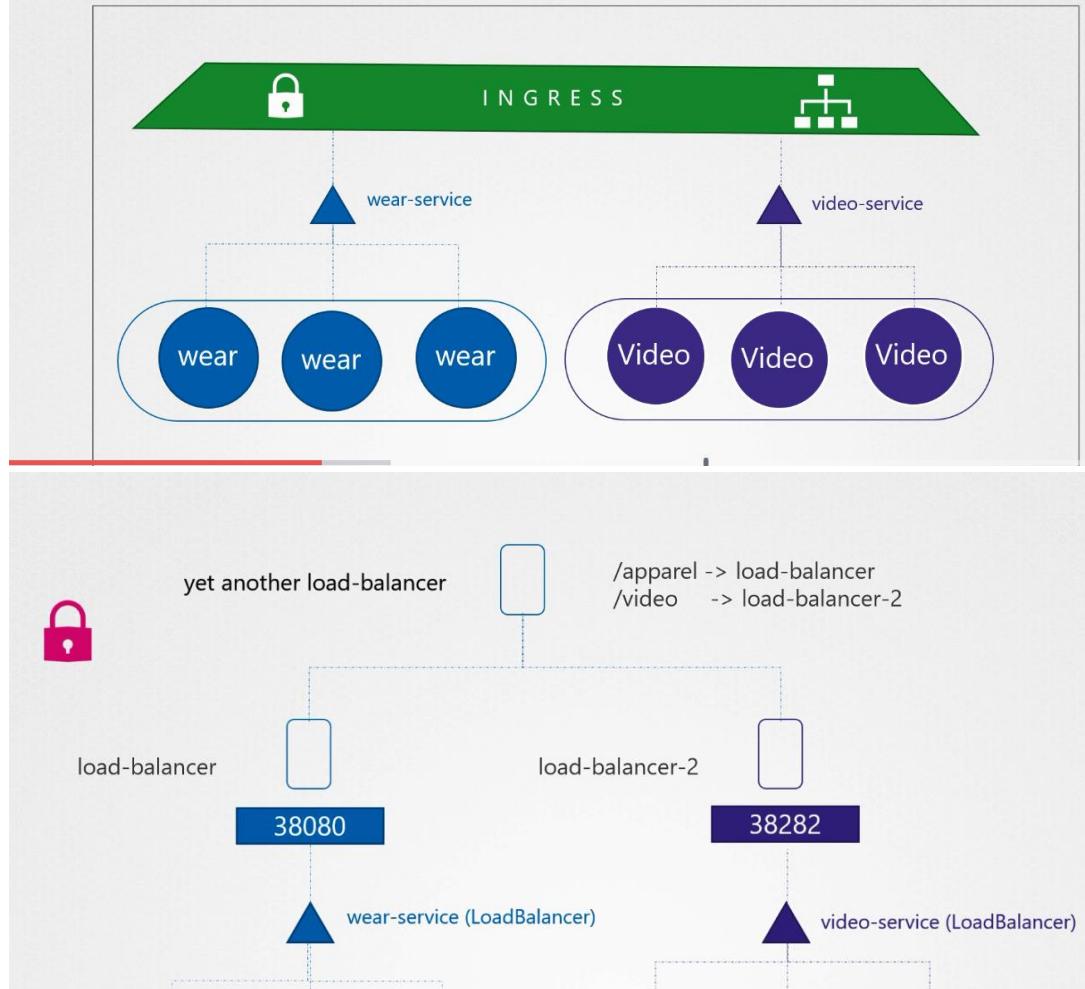
```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: myapp
    type: front-end
```

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: back-end
spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80
  selector:
    app: myapp
    type: back-end
```



Ingress



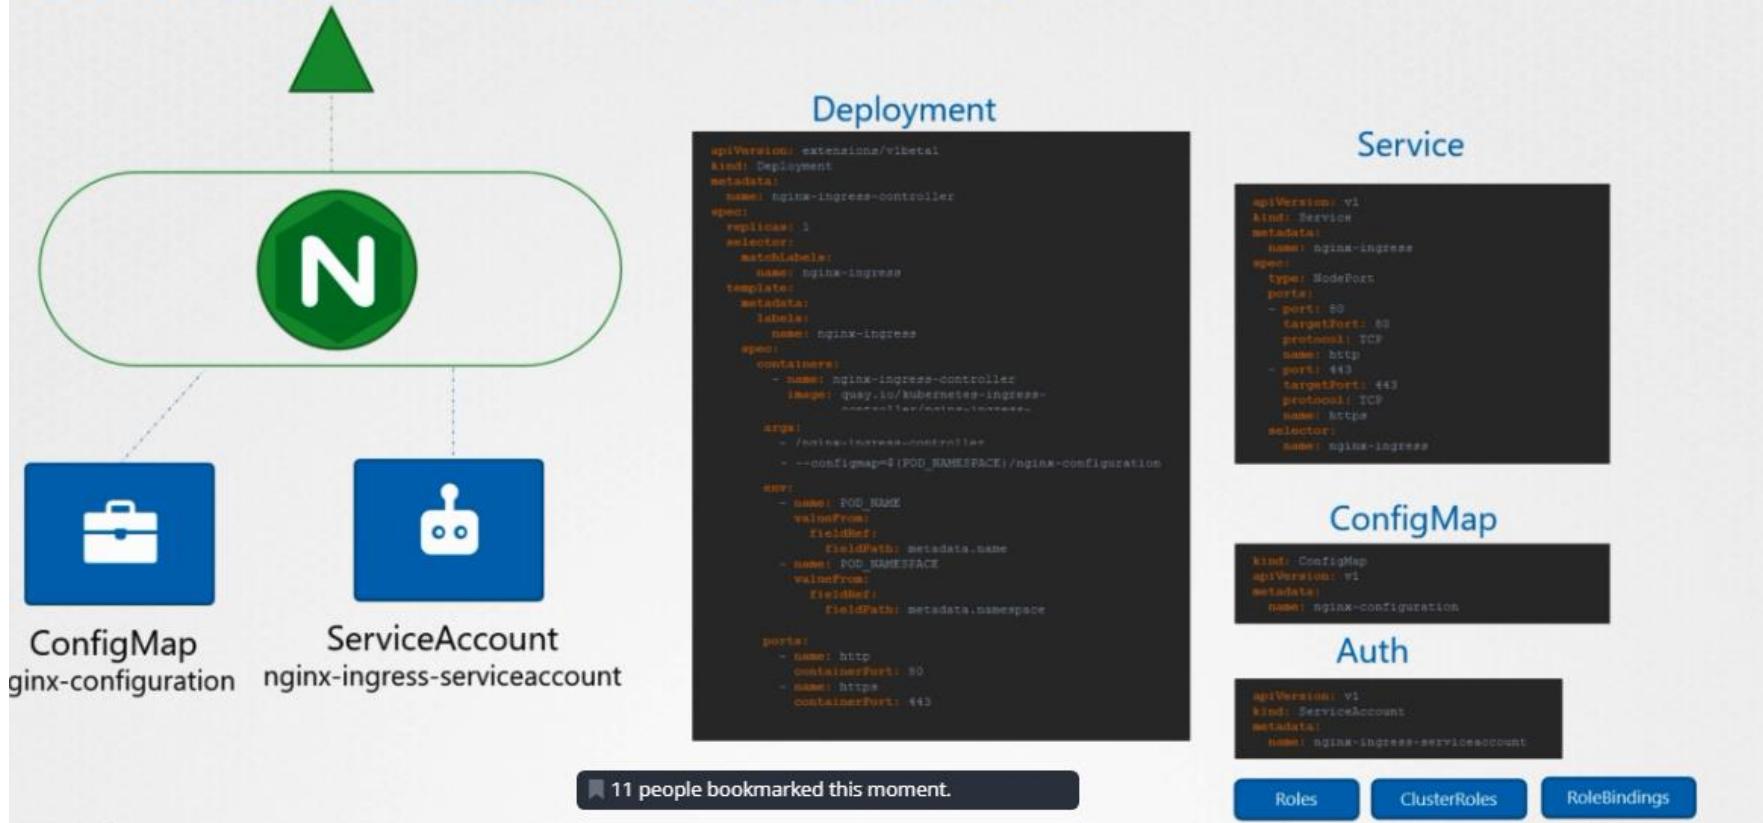
Ingress



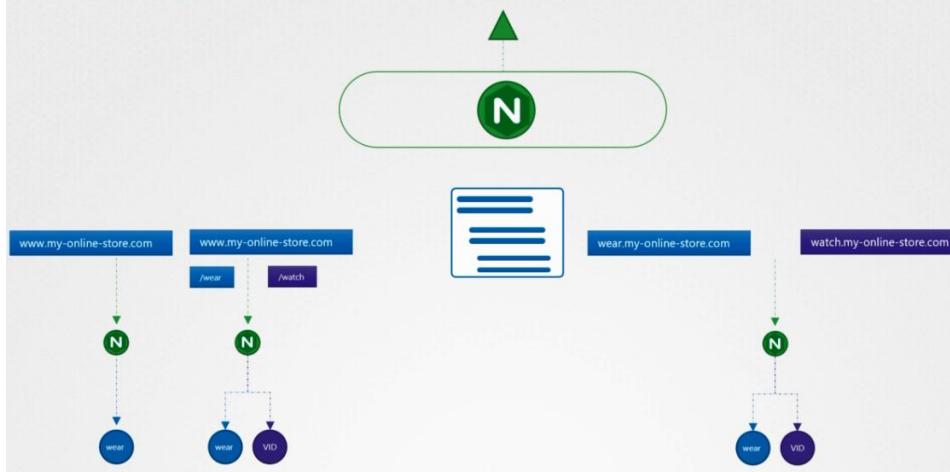
INGRESS CONTROLLER



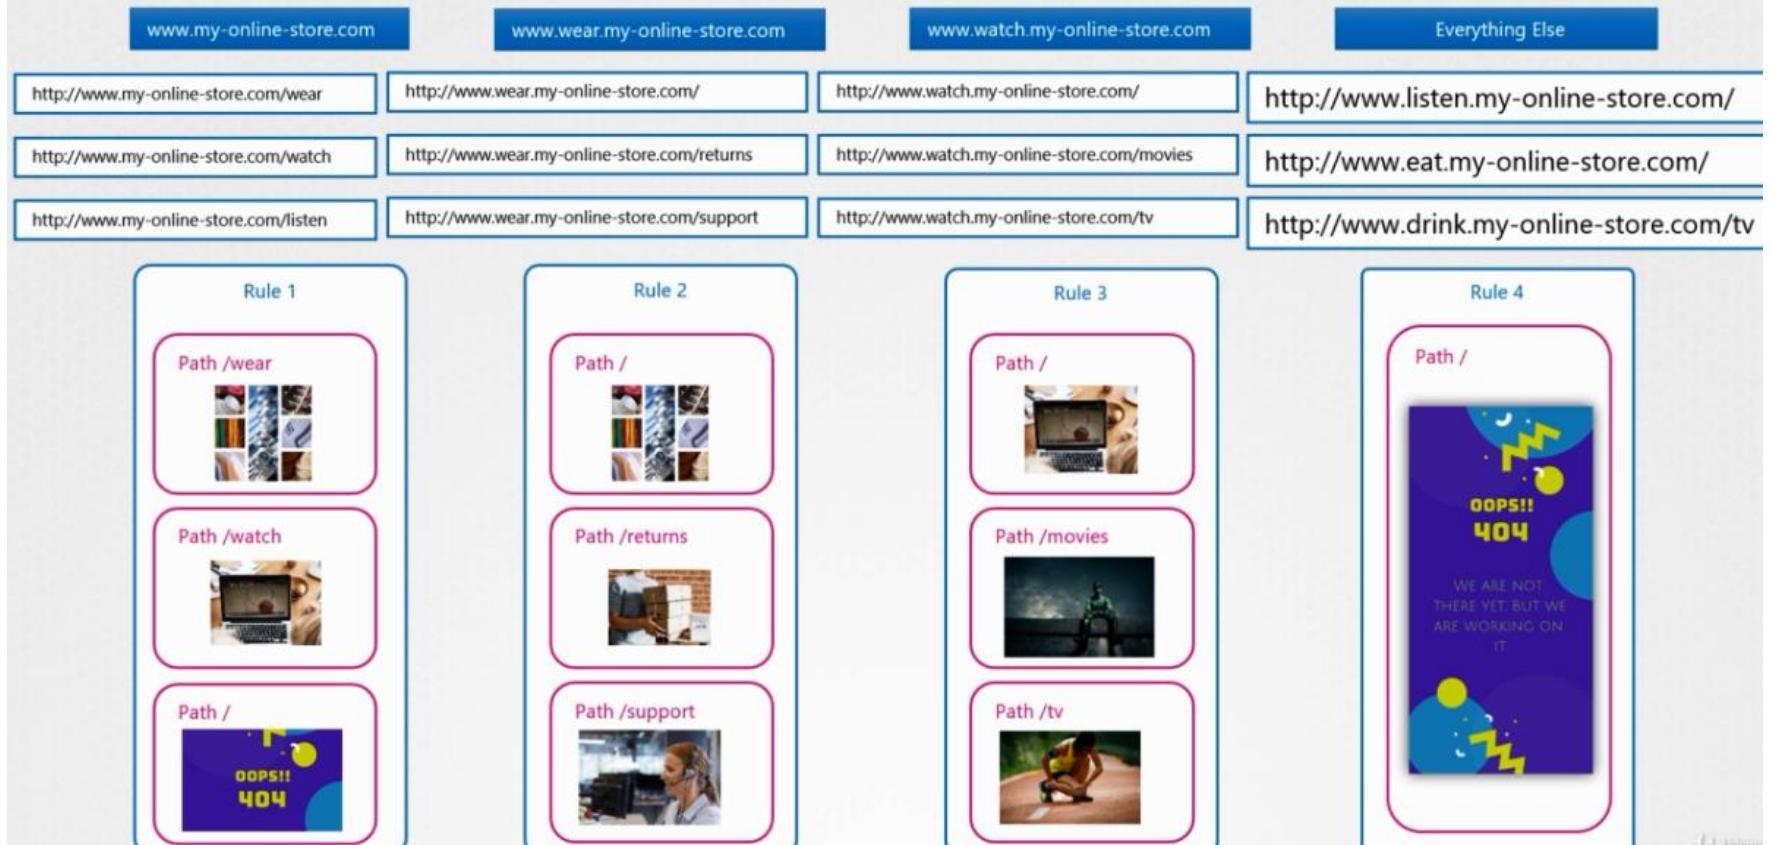
INGRESS CONTROLLER



INGRESS RESOURCE



INGRESS RESOURCE - RULES



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-wear-watch
spec:
  rules:
  - http:
    paths:
    - path: /wear
      backend:
        serviceName: wear-service
        servicePort: 80
    - path: /watch
      backend:
        serviceName: watch-service
        servicePort: 80
```

INGRESS RESOURCE

2 Rules



```
Ingress-wear-watch.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-wear-watch
spec:
  rules:
    - host: wear.my-online-store.com
      http:
        paths:
          - backend:
              serviceName: wear-service
              servicePort: 80
    - host: watch.my-online-store.com
      http:
        paths:
          - backend:
              serviceName: watch-service
              servicePort: 80
```

INGRESS RESOURCE

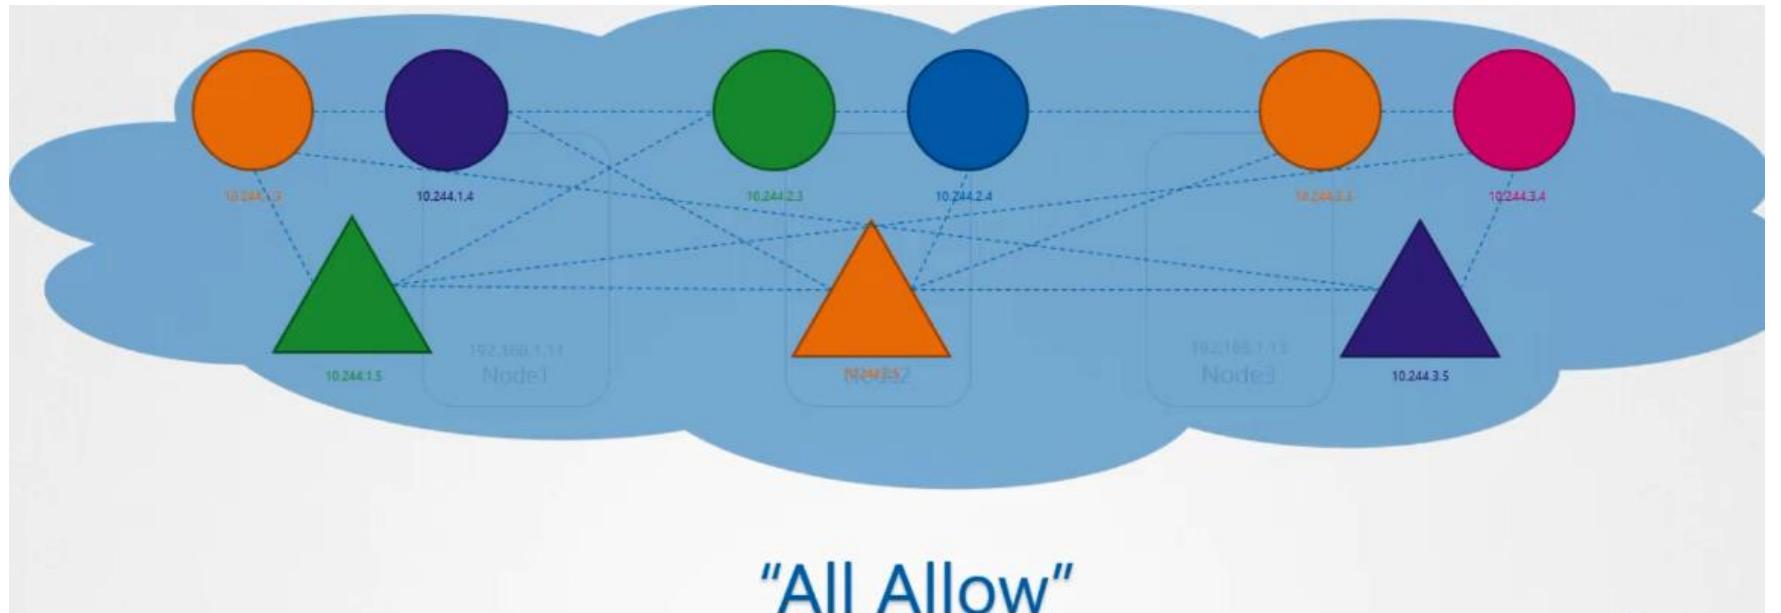
Ingress-wear-watch.yaml

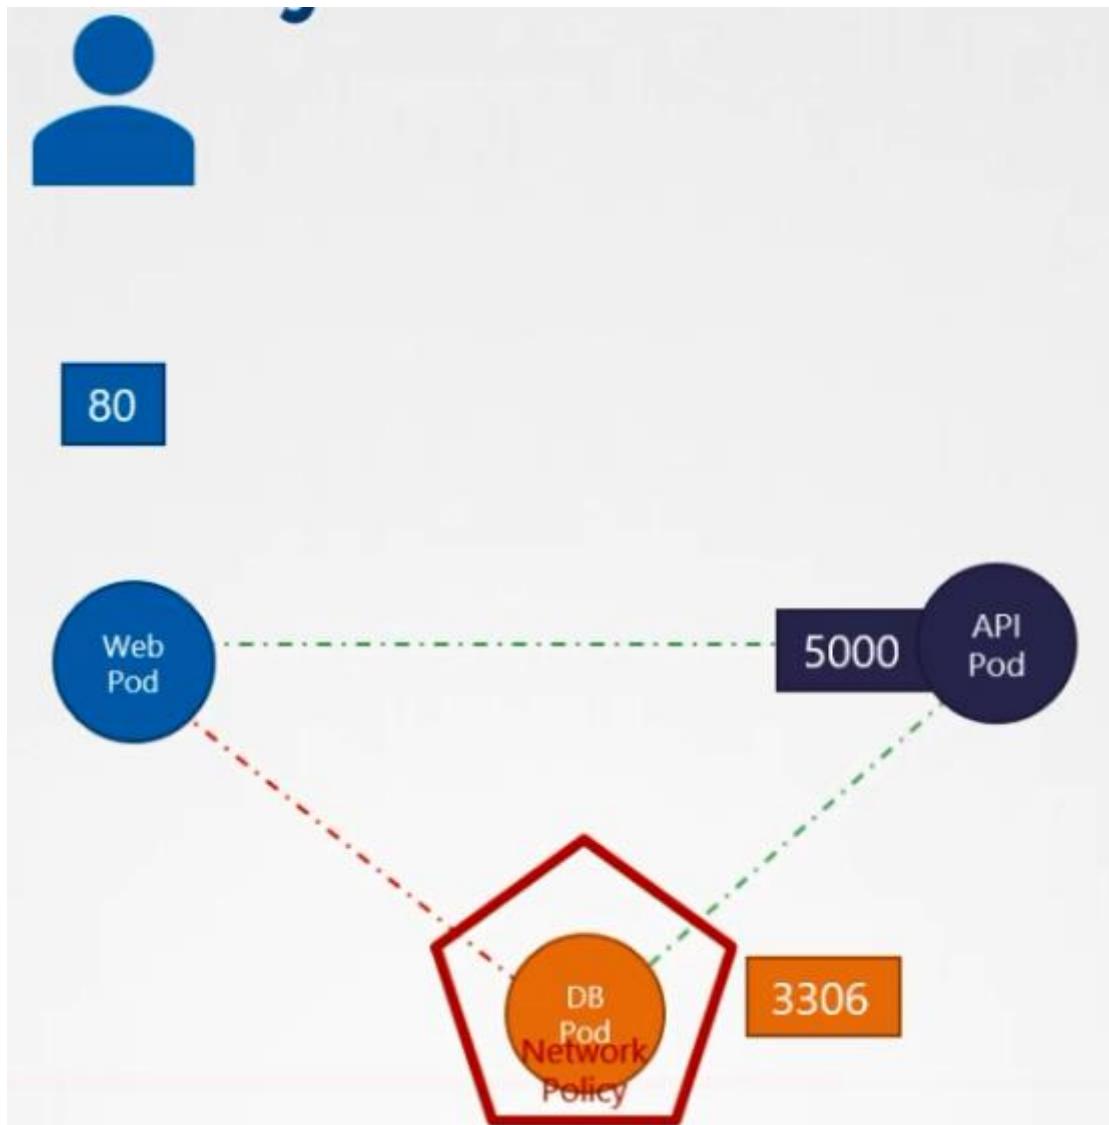
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-wear-watch
spec:
  rules:
  - http:
    paths:
    - path: /wear
      backend:
        serviceName: wear-service
        servicePort: 80
    - path: /watch
      backend:
        serviceName: watch-service
        servicePort: 80
```

Ingress-wear-watch.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-wear-watch
spec:
  rules:
  - host: wear.my-online-store.com
    http:
      paths:
      - backend:
          serviceName: wear-service
          servicePort: 80
  - host: watch.my-online-store.com
    http:
      paths:
      - backend:
          serviceName: watch-service
          servicePort: 80
```

By Default each pod can comminicate with all pod. But some cases it should not be , that's where we specify network policy where we set rules which pod can communicat with which pod





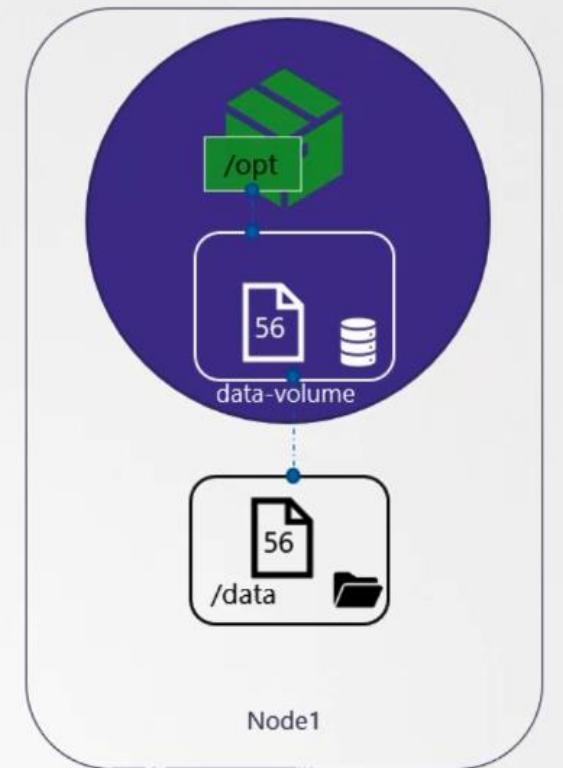
Network Policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: db-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          name: api-pod
  ports:
  - protocol: TCP
    port: 3306
```

State Persistence – 8%

Volumes & Mounts

```
apiVersion: v1
kind: Pod
metadata:
  name: random-number-generator
spec:
  containers:
    - image: alpine
      name: alpine
      command: ["/bin/sh","-c"]
      args: ["shuf -i 0-100 -n 1 >> /opt/number.out;"]
  volumeMounts:
    - mountPath: /opt
      name: data-volume
  volumes:
    - name: data-volume
      hostPath:
        path: /data
        type: Directory
```



```
volumes:  
- name: data-volume  
  awsElasticBlockStore:  
    volumeID: <volume-id>  
    fsType: ext4
```



Persistent Volume

pv-definition.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-voll
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 1Gi
  hostPath:
    path: /tmp/data
```

pv-definition.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-voll
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 1Gi
  awsElasticBlockStore:
    volumeID: <volume-id>
    fsType: ext4
```

Persistent Volume Claim

pvc-definition.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

pv-definition.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-vol1
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 1Gi
  awsElasticBlockStore:
    volumeID: <volume-id>
    fsType: ext4
```

▶ kubectl get persistentvolumeclaim

| NAME | STATUS | VOLUME | CAPACITY | ACCESS MODES | STORAGECLASS | AGE |
|---------|--------|---------|----------|--------------|--------------|-----|
| myclaim | Bound | pv-vol1 | 1Gi | RWO | | 43m |

```
▶ kubectl delete persistentvolumeclaim myclaim  
persistentvolumeclaim "myclaim" deleted
```

persistentVolumeReclaimPolicy: Retain

persistentVolumeReclaimPolicy: Recycle

persistentVolumeReclaimPolicy: Delete

Use Shortcuts/Aliases

po for PODs

rs for ReplicaSets

deploy for Deployments

svc for Services

ns for Namespaces

netpol for Network policies

pv for Persistent Volumes

pvc for PersistentVolumeClaims

sa for service accounts